

# Was magnitude estimation necessary?

A secondary-data reassessment of Bard, Robertson, and Sorace (1996)

Brett Reynolds \*

Humber Polytechnic & University of Toronto

June 2026

## Abstract

This paper reassesses Bard, Robertson, and Sorace’s case for magnitude estimation in linguistic acceptability research. It separates two claims that are often run together: acceptability judgments are graded, measurable, and methodologically serious; magnitude estimation has a practical advantage over simpler formal judgment tasks. The first claim is the stronger legacy. The second requires later method-comparison evidence. The paper sets up a constrained, pre-specified secondary analysis of Sprouse and colleagues’ public method-comparison data, gated on data-reuse permission, uses Langsford et al. as article-level evidence about reliability and response-style risks, and keeps computational benchmarks separate as afterlife evidence rather than evidence about magnitude estimation itself.

Keywords: magnitude estimation; acceptability judgments; grammaticality judgments; secondary data; experimental syntax

## I THE PROBLEM

Bard, Robertson, and Sorace made a methodological claim that was larger than a single elicitation format. Acceptability judgments, they argued, could be treated as graded measurements rather than as informal binary verdicts (Bard et al., 1996). That move helped make experimental syntax a measurement enterprise: judgments could be collected from naive participants, scaled, compared across groups, and subjected to ordinary statistical tests.

The harder question is whether magnitude estimation was necessary for that advance. Magnitude estimation promised a way around the scale restrictions of traditional judgment tasks, giving participants an open response scale and preserving fine distinctions. But later acceptability work did not settle into one privileged elicitation method. Likert ratings, forced choice, yes/no tasks, and Thurstonian models all became part of the measurement toolkit (Langsford et al., 2018; Sprouse & Almeida, 2017; Sprouse et al., 2013).

This paper separates two claims that are often allowed to travel together. One claim is strong and still plausible: acceptability is graded, measurable, and methodologically serious. The other is

---

\*Contact: [brett.reynolds@humber.ca](mailto:brett.reynolds@humber.ca)

weaker: magnitude estimation has a practical measurement advantage over simpler formal tasks. The first claim can survive even if the second does not.

The empirical half of this is not untouched. The methodology literature has already compared elicitation formats and largely deflated magnitude estimation's special status. Schütze (2016) reports that most or all of its advertised advantages over Likert tasks have since been refuted, pointing to Weskott and Fanselow (2011) and Sprouse et al. (2013); later work weighed the formats on convergence (Sprouse et al., 2013), statistical power (Sprouse & Almeida, 2017), and reliability and response style (Langsford et al., 2018).

So the bare verdict that magnitude estimation carries no special advantage is not what this paper adds. The contribution is to separate that verdict from the durable measurement claim, to relocate acceptability within a projectibility account, and to route each later dataset to a prespecified estimand, so that any reanalysis of public comparison data, once reuse is permitted, cannot quietly select whichever scale looks best.

The reassessment is therefore not a rejection of Bard et al.'s programme. It is a test of what should be credited to magnitude estimation itself, and what should instead be credited to the broader formalization of acceptability judgment methodology.

## 2 THE BARD TARGET

The target has to be reconstructed before it is reassessed. Bard et al. present magnitude estimation as a solution to measurement-scale problems in conventional acceptability elicitation, with enough resolution to capture fine-grained linguistic distinctions and enough robustness to work with linguistically naive participants (Bard et al., 1996). That is the source-grounded claim this paper can evaluate against later method-comparison evidence.

A stronger claim sometimes follows in reception: magnitude estimation is not just one formal method among others, but the method with privileged measurement status. That stronger reading should not be attributed to Bard et al. without textual support. It is better treated as a separable claim about what parts of the field took magnitude estimation to have shown.

The distinction matters for fairness and for design. If Bard et al. are read as saying only that acceptability judgments need a graded measurement method, then later Likert, forced-choice, yes/no, and Thurstonian work may extend their programme rather than refute it. If magnitude estimation is read as having a method-specific privilege, then later comparisons have to ask whether it outperforms other tasks on prediction, reliability, decision stability, or construct interpretation.

The available evidence also limits the claim. No public Bard participant-level rows have been found, so this paper cannot directly replicate the original experiment. The empirical reassessment has to be a conceptual replication and secondary-data test: later datasets are used to ask whether the method-specific part of the Bard claim retains independent evidential force.

## 3 SECONDARY-DATA DESIGN

The secondary-data design treats the project as a measurement reassessment of Bard, Robertson, and Sorace (Bard et al., 1996). Four evidential routes should stay separate. Published summaries from Bard et al. can source the historical claim, but not a direct replication unless their participant-level data are found.

Later judgment studies can support secondary analysis when they compare elicitation methods

| Source                | Evidence level                     | Use in this paper                                                 | Claim it cannot support                                        |
|-----------------------|------------------------------------|-------------------------------------------------------------------|----------------------------------------------------------------|
| Bard et al. 1996      | Published article                  | Historical target and scale-resolution claim                      | Direct replication without participant rows                    |
| Sprouse 2013/2017     | Public row-level comparison data   | Constrained secondary method comparison, pending reuse permission | Ungated publication-ready novel results while reuse is pending |
| Langsford et al. 2018 | Published article and supplement   | Reliability and response-style synthesis                          | New participant-level reliability model without raw rows       |
| CoLA, BLiMP, SAD      | Benchmark labels and documentation | Afterlife of formal acceptability measurement                     | Evidence about ME versus other human judgment methods          |

Table 1: Evidential routing before analysis. ME = magnitude estimation; SAD = Syntactic Acceptability Dataset.

on shared or comparable items and provide the rows needed for the intended estimand. Computational benchmarks can show the afterlife of acceptability labels, but they do not bear on magnitude estimation versus Likert ratings. Construct-level discussion can separate acceptability from grammaticality, naturalness, standardness, and correction.

This creates a genre gate. Public row-level data from the Sprouse comparison studies can support a constrained secondary analysis of method performance. The Langsford et al. reliability study is essential evidence for the argument, but unless its raw participant responses are obtained it can only constrain the paper at the article-summary level. The paper should therefore not promise a full response-style or test-retest model across the whole literature.

The Sprouse files also require representation choices before modelling. For Sprouse et al. (2013), the primary forced-choice comparison should use the pair-level sign-test file, because forced choice is a pairwise task and that file already records judgments against expected patterns. The logistic file preserves item-level binary rows, changing the unit of analysis, so it belongs as a named sensitivity specification.

For Sprouse and Almeida (2017), yes/no rows with item IDs B, G, and M should be excluded from item-level comparisons. They are absent from the materials workbook, occur in repeated early order positions for every subject, and do not align with the magnitude-estimation, Likert, or forced-choice item universe. This exclusion is a structural inclusion rule, not an outcome-based trimming decision.

Before any modelling, each source should be assigned to an estimand it can actually answer. The primary estimand is the item-level acceptability signal recoverable across methods. Secondary estimands are method-specific noise or bias and decision agreement for theoretically relevant contrasts.

Participant response style and method-by-item instability are admissible only when the dataset contains the participant-level rows and repeated structure needed to estimate them. This ordering keeps the analysis from selecting whichever scale looks best after the fact, the kind of forking path Gelman and Loken warn about (Gelman & Loken, 2013, 2014).

The substantive comparison is a common-latent-signal null: magnitude estimation, Likert rat-

ings, forced choice, and Thurstonian choice are first modelled as noisy observations of the same item-level acceptability. A magnitude-estimation term becomes interesting only if it improves prediction, reliability, or decision stability enough to matter. Scrambled labels, method-label permutations, and simulations from a no-method-advantage latent model are negative controls for the pipeline, checking whether the analysis manufactures a method winner from degraded structure. They are not substitutes for the substantive common-signal comparison.

#### 4 SCALE AND RELIABILITY

The scale question should be phrased as practical measurement advantage over the common-latent null. Bard et al.'s case for magnitude estimation included scale resolution, fine-grained distinctions, naive-participant use, and replication across groups (Bard et al., 1996). The safer target is that source-grounded claim, together with the stronger reading of magnitude estimation that took hold in parts of the later methodology literature. Later method-comparison data give the paper a sharper question: does magnitude estimation recover item differences, participant agreement, or theoretically relevant decisions that simpler tasks lose?

Sprouse et al. (2013) compare formal judgment tasks with informal published judgments, giving evidence about convergence between magnitude estimation, Likert ratings, and forced choice. The public Sprouse rows can support a bounded reanalysis of item-level agreement and decision consequences, subject to reuse permissions. Langsford et al. (2018) add a reliability and response-bias target by comparing Likert, magnitude estimation, forced choice, and Thurstonian models. Without their raw rows, Langsford can support the synthesis of published reliability findings but not a new participant-level response-style or test-retest model.

The reliability analysis inherits the unit-of-analysis charter fixed in the secondary-data design rather than reopening it. The same primary and sensitivity representations carry over, so reliability claims stay traceable to fixed units of analysis instead of to whichever representation later proves most convenient.

The primary reliability analysis should compare methods on prediction, stability, decision consequences, mean separation, and correlation, with priority on consequences for the linguistic contrast. A useful result would show whether magnitude estimation changes the answer for the linguistic contrast under study, reduces Type S or Type M risks (Gelman & Carlin, 2014), or carries costs that offset any added resolution. Sensitivity checks should vary transformations and exclusion rules as named specifications rather than as invisible post hoc choices (Steege et al., 2016).

#### 5 THE CONTEMPORARY AFTERLIFE OF ACCEPTABILITY

Contemporary benchmark resources matter because they show what survived from the formal acceptability programme. CoLA turns examples from linguistics publications into binary labels for model evaluation (Warstadt et al., 2019). BLiMP turns expert-designed grammatical contrasts into large sets of minimal pairs (Warstadt et al., 2020). These resources inherit the idea that acceptability-related contrasts can be made portable, scored, and compared across systems.

They also change the evidential object. CoLA's labels come from published linguistic examples rather than from a new sample of naive participant responses. BLiMP's contrasts are generated from expert-crafted grammars and then used to test whether a model prefers the acceptable sentence in a pair. Those are useful evaluation resources, but they are not method-comparison data for magnitude

estimation, Likert ratings, forced choice, or yes/no judgments.

The Syntactic Acceptability Dataset is the most useful bridge case because it keeps grammatical status and crowd acceptability status apart (Juzek, 2024). That separation is exactly the level discipline this paper needs. Benchmark labels can show the afterlife of formal acceptability measurement, while participant-response datasets are needed to compare human elicitation methods.

## 6 GRAMMATICALITY AND ACCEPTABILITY

The reassessment also needs a disciplined vocabulary. The background view adopted here is that grammar is a system of projectibility profiles (Goodman, 1955): relatively stable patterns that license inferences from forms, meanings, contexts, speakers, and judgments to uses a community or register treats as available. A grammaticality claim is therefore not a report of one response. It states what a form projects to: interpretation, production, correction, expert classification, register fit, cross-speaker stability, and related consequences.

Participant responses under an acceptability prompt are observations of judgments under a task. Expert asterisks and published example labels are classifications made within a theoretical or editorial practice. Benchmark labels are dataset-construction decisions. Each can enter a grammaticality profile, but none is identical to the profile.

This discipline keeps neighbouring constructs apart. Acceptability is one evidential route into a projectibility profile, not the profile itself. Personal production, naturalness, standardness, correction behaviour, and theoretical classification project differently. A sentence can be dispreferred without being ungrammatical, correctable without being unnatural, or standard in one register while unacceptable under another task framing.

Bard et al. helped make one part of this space measurable: graded acceptability judgments from participants (Bard et al., 1996). That is a major contribution to the measurement of one evidential channel. It does not by itself show that an acceptability scale captures every inference licensed by a grammar. Later benchmarks sometimes repackage acceptability-related labels as grammaticality targets (Warstadt et al., 2019, 2020). The Syntactic Acceptability Dataset is valuable here because it explicitly separates literature-derived grammatical status from crowd acceptability status (Juzek, 2024).

That boundary also limits the no-new-data paper. Existing datasets can compare the stability of elicitation methods and clarify how acceptability labels travel. They cannot settle a full projectibility profile for a construction: production, correction, register, interpretation, and community licensing would require data designed for those inferences.

## 7 CONCLUSION

The balanced update is straightforward. Bard et al. were right to make acceptability measurable and methodologically serious. Their documented case for magnitude estimation concerned scale resolution, fine distinctions, naive-participant use, and replication across groups. That claim deserves credit as part of the history of experimental syntax.

The method-specific conclusion is more limited. Magnitude estimation has to be compared against a common-latent-signal null, where later formal tasks are treated first as noisy measures of the same item-level acceptability signal. The question is whether magnitude estimation improves prediction, reliability, decision stability, or construct interpretation enough to matter for linguistic work.

The paper's final genre depends on the data gate. If Sprouse reuse permission is available, the project can become a constrained secondary analysis of row-level method-comparison data. If reuse remains unavailable, the project should stay a synthesis and design paper, using Sprouse for verified public structure and Langsford et al. for article-level reliability evidence.

In either genre, the central conclusion is the same. The durable legacy is the formal measurement of graded acceptability as one route into grammar's projectibility profiles. The privileged status of magnitude estimation is a separate claim, and later evidence has to earn it separately.

#### DATA AND CODE AVAILABILITY

No raw participant-level data are committed to this repository. Source and dataset verification are tracked in `notes/source-verification.md`; the analysis charter and structural Sprouse inventory are tracked in `notes/analysis-charter.md` and `notes/dataset-inventory.md`. Scripts currently generate ignored local structural derivatives and readiness manifests, while novel outcome analysis remains gated pending Sprouse data reuse permission or an explicitly verification-only framing.

#### ACKNOWLEDGEMENTS

TODO: Add final AI-use disclosure with the actual model list used on the project.

## REFERENCES

- Bard, E. G., Robertson, D., & Sorace, A. (1996). Magnitude estimation of linguistic acceptability. *Language*, 72(1), 32–68. <https://doi.org/10.2307/416793>
- Gelman, A., & Carlin, J. (2014). Beyond power calculations: Assessing type S (sign) and type M (magnitude) errors. *Perspectives on Psychological Science*, 9(6), 641–651. <https://doi.org/10.1177/1745691614551642>
- Gelman, A., & Loken, E. (2013). *The garden of forking paths: Why multiple comparisons can be a problem, even when there is no “fishing expedition” or “p-hacking” and the research hypothesis was posited ahead of time* [Unpublished manuscript, Columbia University]. [https://sites.stat.columbia.edu/gelman/research/unpublished/p\\_hacking.pdf](https://sites.stat.columbia.edu/gelman/research/unpublished/p_hacking.pdf)
- Gelman, A., & Loken, E. (2014). The statistical crisis in science. *American Scientist*, 102(6), 460–465. <https://doi.org/10.1511/2014.111.460>
- Goodman, N. (1955). *Fact, fiction, and forecast*. Harvard University Press.
- Juzek, T. S. (2024). The Syntactic Acceptability Dataset (preview): A resource for machine learning and linguistic analysis of English. *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, 16113–16120.
- Langsford, S., Perfors, A., Hendrickson, A. T., Kennedy, L. A., & Navarro, D. J. (2018). Quantifying sentence acceptability measures: Reliability, bias, and variability. *Glossa: a journal of general linguistics*, 3(1), Article 37, 1–34. <https://doi.org/10.5334/gigl.396>
- Schütze, C. T. (2016). *The empirical base of linguistics: Grammaticality judgments and linguistic methodology*. Language Science Press. <https://doi.org/10.17169/langsci.b89.100>
- Sprouse, J., & Almeida, D. (2017). Design sensitivity and statistical power in acceptability judgment experiments. *Glossa: A Journal of General Linguistics*, 2(1), 1–32. <https://doi.org/10.5334/gigl.236>
- Sprouse, J., Schütze, C. T., & Almeida, D. (2013). A comparison of informal and formal acceptability judgments using a random sample from *Linguistic Inquiry* 2001–2010. *Lingua*, 134, 219–248. <https://doi.org/10.1016/j.lingua.2013.07.002>
- Steege, S., Tuerlinckx, F., Gelman, A., & Vanpaemel, W. (2016). Increasing transparency through a multiverse analysis. *Perspectives on Psychological Science*, 11(5), 702–712. <https://doi.org/10.1177/1745691616658637>
- Warstadt, A., Parrish, A., Liu, H., Mohananey, A., Peng, W., Wang, S., & Bowman, S. R. (2020). BLiMP: The benchmark of linguistic minimal pairs for English. *Transactions of the Association for Computational Linguistics*, 8, 377–392. [https://doi.org/10.1162/tacl\\_a\\_00321](https://doi.org/10.1162/tacl_a_00321)
- Warstadt, A., Singh, A., & Bowman, S. R. (2019). Neural network acceptability judgments. *Transactions of the Association for Computational Linguistics*, 7, 625–641. [https://doi.org/10.1162/tacl\\_a\\_00290](https://doi.org/10.1162/tacl_a_00290)
- Weskott, T., & Fanselow, G. (2011). On the informativity of different measures of linguistic acceptability. *Language*, 87, 249–273.